

Horus cheatsheet

This document will contain checklists to help be sure your task meets all guidelines.

Does the conversation meet requirements?

- ☐ There is at least **one failure** (the model outputting incorrect information to the user)
- ☐ **Five tool calls** are present in each turn
- ☐ **Five servers** are represented across all task turns
- ☐ The task is written from the point of view of the correct **persona, domain, and lens**
- ☐ Prompts sound natural, like believable requests that a person would make of AI. They are context-rich using information from the provided persona.
- ☐ The turns together are a believable, interconnected conversation, with carried context

Is the Golden Trajectory the “ideal agent trajectory”?

- ☐ Are all parameters and output information **factually correct / grounded**?
- ☐ There are **no redundant tool calls**
- ☐ **All output information is grounded** in outputs from tool calls
- ☐ All user-provided information is verified
- ☐ The approach taken for each turn’s trajectory effectively addresses the user’s intent
- ☐ The Golden Trajectory addresses the all user intentions from all requests
- ☐ Limit parameter choices are consistent with addressing the user’s intention
- ☐ If checking the date, are you using the **brave_search_web_search** tool?

Are the rubrics compliant?

- ☐ Rubric comprehensively covers all critical tool calls, information identification/carryover/manipulation from prior turns in multi-step planning rubrics
- ☐ If a trajectory passes all rubric criteria, that trajectory would be guaranteed to produce the correct outputs at each step
- ☐ There are **no negative rubrics** (all rubrics check for the presence of the correct action/output not the absence of some action/output)
- ☐ Each rubric criterion is correctly labeled with:
 - Tool selection (for choosing each unique tool used)
 - Query construction (for passing the correct parameters to tools)
 - Multi-step planning (for identifying or manipulating information)
 - Output (covering correctness and completeness of all output information)
- ☐ Rubric special cases are handled correctly, to avoid over-specifying:
 - Query Construction rubric criteria for “fuzzy search”** tools with this format: “the model must set the ‘query’ parameter to clearly target [target] (e.g., X, Y) when calling the [specific tool]”
 - Tool Selection / Query Construction / Multi-step Planning Criteria for Different tool calls satisfying the user’s intent:** “The model must use A tool or an enabled tool, such as B” (check other examples in Rubrics Writing Guidelines)
- ☐ Rubrics are atomic, self-contained, and non-overlapping.