

Frequently Asked Questions

Q: Does the model have to include 5 turns per “assistant turn” within the Agent and Golden Trajectory?

A: No. The model needs to include 5 tool calls across the entire turn, not 5 for each assistant turn within the trajectory.

🔴 **Note:** “Assistant turn” and “turn” are not the same. A turn can include one or more assistant turns.. Each assistant turn may include one or more tool calls, depending on how the actions to fulfill the prompt need to be split.

Q: Do the five different servers have to be used in every turn?

A: No. The task needs to include five different servers, but they can be distributed across turns as needed.

Q: What is the “Agent Environment” (the section that appears before writing each prompt) used for?

A: The Agent Environment plays a key role in prompt creation and in setting the model’s context. It’s where we interact with tools to keep prompts realistic and natural, and where we get context on how the model is expected to behave (for example, which tools to use and what parameters to pass to get results).

Q: How do I create natural prompts that trigger 5+ tool calls?

A: By putting yourself in the user’s (persona’s) shoes and thinking about how the model could realistically help solve situations or problems within that context.

Q: Are redundant tool calls considered a “Heavy Failure” (model stump)?

A: No. Redundant or refinement tool calls should be corrected by keeping only the tool that most efficiently retrieves the results, but the model should not be failed for this. A score of 3 or 4 is sufficient penalization in these cases.

Q: How should we deal with the ‘limit’ parameter that is being used in many tools?

A: We should consider whether using the default value or a different value for the ‘limit’ parameter has a real impact on retrieving results that are relevant to the user’s intent. It is not the same when a user wants results from the entire perfume shop and the model applies a limit of 50, leaving 150 products out of the tool output, versus cases where the specific limit value of 10, 20, or 50 does not meaningfully affect the results retrieved.

This is one of the reasons why interacting with the Agent Environment is so important, to check whether the applied limit, according to the user’s request, is acting as a blocker or is simply irrelevant.

Q: Do I have to add the ‘limit’ parameter as a Query Construction rubric criteria?

A: It will depend on whether it is relevant to the prompt’s fulfillment, as mentioned above. In a lot of cases, the answer will be no.

Q: Do I have to include multi-step rubric criteria for every element in the tool output?

A: Only for the information that plays a role in the final user facing response and related to the user intent. If the prompt asks for products name and price, the rubric must include identification (multi-step) criteria for those two pieces of information. This can be a single rubric covering both, or two separate rubrics (recommended), one for each.

Q: If the task specifies a maximum of 6, 7, or 8 turns, does that mean I have to include that number of turns in the trajectory?

A: No. You only need to include a minimum of 5 turns.

Q: How can I meet the enterprise lens in a persona that is not natural to be a business man, like Jordan Williams (teenager)?

A: The enterprise lens is not limited to business ownership. It can be expressed through leadership or management responsibilities that align with the persona, such as leading school projects, organizing student groups, managing sports teams, coordinating events, or owning responsibility for budgets, planning, and execution.