Errors that still haven't been addressed:

1. Prompts are still sometimes under-specified.
   If your prompt does not make it clear how the final answer should be selected, the task becomes hard to solve deterministically and hard to rubric correctly. When a search returns many valid results (you should know this before crafting your prompts by using the playground), provide a clear selection rule. If that rule can tie, add a tie-breaker.
2. Cross-turn continuity still needs to be stronger.
   All the turns should truly depend on earlier information or tool results. Staying on the same topic is not enough on its own. Later turns should build on known constants established earlier in the conversation.
3. Golden trajectories still sometimes drift away from the tool output.
   The tool output is the source of truth. Do not add facts that are not grounded in the trajectory unless they are clearly common knowledge and are presented that way.
4. Some GTs still include unnecessary calls or weak search logic.
   The ideal path should be efficient and fully justified. Avoid redundant retrieval, unnecessary refinement, or extra steps that do not help resolve the request. Double-check that your new tool call hasn't been made in previous turns.
5. Rubrics still miss important coverage or become too prescriptive.
   Rubrics must cover the real must-haves of the turn, both process and outcome. At the same time, do not overprescribe a single exact path when multiple valid paths exist. Query-construction over-specificity is still something to watch closely. Please do not mirror the Golden Trajectory. Think about all the possibilities.
6. Assistant messages still need attention.
   Every assistant turn needs a message explaining the action's intent. The message should not claim an outcome before the tool runs, and it should stay coherent with the actual tool call(s).

New things to keep an eye on:

- 5-tool-call compliance is now being flagged more often.
  A recent pattern in reviewer feedback is that the turns do not meet the minimum tool-call complexity requirement.
  - Each turn should include at least 5 tool calls, and each task should use at least 5 distinct servers.
  - Build prompts that naturally require that level of complexity rather than trying to force it late.
  - The tool calls must be useful and meaningful.

    - **Do not treat simple calculator use as a justification to increase the count**. If the calculation is trivial and does not add real value to solving the turn, it should not be used to pad tool-call complexity.
    - The goal is meaningful tool use, not artificially reaching the minimum.
- Time grounding needs to be tighter.
  If the prompt depends on time, especially relative windows like “last 14 days,” “first week of February,” or any case with an unstated year, make sure the trajectory verifies the exact date or timestamps instead of assuming. Time-variant prompts must be grounded explicitly.
- Be careful with assistant-turn sequencing.
  Do not place dependent tool calls in parallel in the same assistant turn. If one call depends on information returned by another, it must come later. Also, make sure the assistant turn reflects the prompt's real sub-intents rather than collapsing everything into one.

What to keep in mind while tasking

Before finalizing a task, do a quick self-check:

- Does each turn clearly require tools?
- Does each turn naturally reach at least 5 useful and meaningful tool calls?
- Am I avoiding trivial calculator calls that do not materially help solve the turn?
- Do at least three turns depend on earlier known constants?
- Is every time-dependent part explicitly grounded?
- Does the GT say only what the tools support?
- Does each assistant turn explain the intent clearly?
- Do the rubrics cover the must-haves without locking the model into one unnecessarily narrow path?

Excess rubrics that artificially constrain the agent

Rubrics must enforce the agent to perform *the minimally required set of actions which guarantee the required outcome*

In every trajectory, the in-task agent will output more information than is required. And you may even include some extra details in your GT (you should avoid as many as possible).

Do not create rubrics that judge whether the model identifies and outputs information more than is required.

Example: reporting inventory during a price lookup

For instance, when looking up a price, reporting the specific amount of the item in stock may be helpful, but it is not minimally necessary. It should therefore not be included in the rubric.

- It may be necessary, depending on the wording, to check whether the number in stock is greater than zero. But it is not necessary to report the exact number in stock.

Missing rubrics

For each tool call, you will likely have 3 to 5 rubric criteria. Here is a common pattern:

1. Multi-step planning rubric identifying information from the user prompt, prior turn, or prior tool call
2. Tool selection rubric to select the tool (if it wasn't used that turn)
3. Query construction rubric to pass the information/parameter(s) into the tool
4. Multi-step planning rubric identifying the relevant output
5. Either:
    a. Repeat the cycle for the new tool call
    b. Output rubric stating the identified information necessary to output to address the user intent

Multi-Step Planning Rubric Criteria Missing

A recurring rubric-writing mistake is failing to include Multi-Step Planning criteria for information that appears in the final answer.

Every relevant output from a tool call must have a Multi-Step Planning criterion to identify it, even if that same information is also checked again in Output Accuracy & Completeness. “Relevant” means information directly tied to the user’s request, not every single field in the tool output, but the fields the model needs to identify and surface in the user-facing answer.

What this means in practice:

- If the user asks for a **price**, the rubric should include an MSP criterion to identify the **price** in the tool's output.
- If the user asks for **availability/stock**, the rubric should include an MSP criterion to identify **availability/stock** from the tool output.
- If the user requests a specific product or item, the rubric should include an MSP criterion to identify the correct **product/item** in the tool output.
- It is **not enough** to only check those facts in Output Accuracy & Completeness.

A common mistake is writing:

- tool selection criteria,
- query construction criteria,
- final output criteria,

But skip the intermediate identification step in the tool results.

Example:

If the final answer says **Coffee Cake**, **$5.99**, and **18 in stock**, the rubric should not only check that those are reported in the final answer. It should also check that the model correctly identified **Coffee Cake**, its **price**, and its **stock** from the search output.

So the rule of thumb is:

If a piece of information from a tool output, from earlier turns, or from user input is relevant to the user’s request and is used in tool input or appears in the user's output, it should have a matching Multi-Step Planning criterion.